



Faculty of Science


# Cryptographic Library for FPGAs using SME

## A Bachelor Project Defense

Jacob Herbst (mrw148) and  
Jonas Flach-Jensen (sjm233)  
Institute of Computer Science (DIKU)

June 18, 2021

Slide 1/102



# Agenda

- Why is this project interesting?
- Presentation of MD5 and AES
- Results
- suggestions for future work




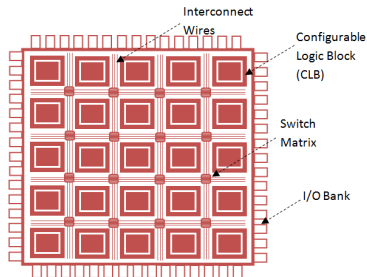
# Motivation - Cryptography

- Is ubiquitous
- Often Hardware-focused
- This project has focused on implementing a variety of algorithms instead of nitpicking one algorithm.



# Motivation - Why use FPGAs?

- Architecture
- Configurable not only in computation but also in interface
- Often fast as the overhead from generality is omitted
- Low power consumption (fixed precision)
- Difficult to program (Xilinx Vivado 🏠)



# MD5

- Cryptographic hash function
- Merkle-Damgaard construction
- Four different compression functions, 64 rounds




Figure: MD5 round




Figure: Merkle-Damgaard construction



# MD5 - Optimizations

- Naive; 1 simple process, 4 busses
- Pipelined version; clocked process for preprocessing and each compression stage
- Same idea for SHA-2




Figure: MD5 pipeline



# MD5 - Pipeline

- Stalls on large messages due to data dependency
- Solution is enabling multiple inputs

Independent message blocks									
clock	0	1	2	3	4	5	6	7	8
	P <sub>1</sub>	M <sub>1</sub>	F <sub>1</sub>	G <sub>1</sub>	H <sub>1</sub>	I <sub>1</sub>	C <sub>1</sub>		
		P <sub>2</sub>	M <sub>2</sub>	F <sub>2</sub>	G <sub>2</sub>	H <sub>2</sub>	I <sub>2</sub>	C <sub>2</sub>	

Dependent message blocks												
clock	0	1	2	3	4	5	6	7	8	9	10	11
	P <sub>1</sub>	M <sub>1</sub>	F <sub>1</sub>	G <sub>1</sub>	H <sub>1</sub>	I <sub>1</sub>	C <sub>1</sub>					
		P <sub>2</sub>	M <sub>2</sub>	-	-	-	-	F <sub>2</sub>	G <sub>2</sub>	H <sub>2</sub>	I <sub>2</sub>	C <sub>2</sub>



# AES

- The algorithm is Rijndael. It is a Block Cipher and a Substitution-permutation (SP) network.
- Four steps, or lookups:

$$W_i \begin{cases} K_i & \text{if } i < N \\ W_{i-N} \oplus \text{SubWords}(W_{i-1} \lll 8) \oplus \text{rcon}_{i/N} & \text{if } i \geq N \text{ and } i \equiv 0 \pmod{N} \\ W_{i-N} \oplus \text{SubWords}(W_{i-1}) & \text{if } i \geq N, N > 6, \text{ and } i \equiv 4 \pmod{N} \\ W_{i-N} \oplus W_{i-1} & \text{otherwise} \end{cases}$$

$$T_0[a] = \begin{bmatrix} S[a] \cdot 02_{16} \\ S[a] \\ S[a] \\ S[a] \cdot 03_{16} \end{bmatrix} \quad T_1[a] = \begin{bmatrix} S[a] \cdot 03_{16} \\ S[a] \cdot 02_{16} \\ S[a] \\ S[a] \end{bmatrix} \quad T_2[a] = \begin{bmatrix} S[a] \\ S[a] \cdot 03_{16} \\ S[a] \cdot 02_{16} \\ S[a] \end{bmatrix} \quad T_3[a] = \begin{bmatrix} S[a] \\ S[a] \\ S[a] \cdot 03_{16} \\ S[a] \cdot 02_{16} \end{bmatrix}$$

$$e_j = T_0[a_{0,3}] \oplus T_1[a_{1,2}] \oplus T_2[a_{2,1}] \oplus T_3[a_{3,0}] \oplus k_j$$





# AES - Optimization

- Fast naive version
- Pipelined by splitting up rounds
- no data dependency



# Results from pipelining

- $hi(x) = x + 2 \cdot blocks$
- $lo(x) = 2 + x \cdot blocks$
- $C(x) = x + 2 \cdot blocks$
- Zedboard bus: 2132 MBps

MD5: Is easily optimised beyond the memory limit.

SHA: Worst hard to optimize because of dependencies.

AES: Reached limit with current approach? Potentially other approaches can reach higher.

ChaCha: Easily reaches limit.

MD5							
Version	$f_{max}$ (Mhz)	$clocks_{hi}$	$TP(MBps)_{hi}$	$clocks_{lo}$	$TP(MBps)_{lo}$	LUT	FF
Naïve	2.38	b	152	b	152	11607	2304
Proc <sub>4</sub>	9.50	hi(6)	266	lo(6)	101	10247	5226
Proc <sub>8</sub>	19.00	hi(10)	532	lo(10)	122	10087	7538
Proc <sub>16</sub>	33.50	hi(18)	937	lo(18)	119	10206	12162
Proc <sub>32</sub>	65.00	hi(34)	1817	lo(34)	123	10149	21347
Proc <sub>64</sub>	<b>115.00</b>	<b>hi(66)</b>	<b>3209</b>	<b>lo(66)</b>	<b>112</b>	<b>10350</b>	<b>39718</b>

SHA							
Version	$f_{max}$ (Mhz)	$clocks_{hi}$	$TP(MBps)_{hi}$	$clocks_{lo}$	$TP(MBps)_{lo}$	LUT	FF
Naïve	2.1	b	134.4	b	134.4	24330	2560
Proc <sub>4</sub>	<b>8.0</b>	<b>hi(6)</b>	<b>223.9</b>	<b>lo(6)</b>	<b>85.3</b>	<b>24466</b>	<b>8938</b>
Proc <sub>8</sub>	8.0	hi(10)	223.8	lo(10)	51.2	24756	14066

AES						
Version	$f_{max}$ (Mhz)	$clocks$	$TP(MBps)$	LUT	FF	BRAM
Naïve	22	b	352	10612	3195	0
TBox	25	b	400	16458	3195	0
Proc <sub>4</sub>	68	C(3)	544	16474	2817	0
Proc <sub>11</sub>	<b>208</b>	<b>C(12)</b>	<b>1663</b>	<b>15659</b>	<b>4383</b>	<b>0</b>
Proc <sub>22</sub>	217	C(24)	1662	15454	7401	0
BRAM <sub>11</sub>	195	C(31)	1556	10012	10398	72

ChaCha					
Version	$f_{max}$ (Mhz)	$clocks$	$TP(MBps)$	LUT	FF
Naïve	1.25	b	80	14670	3457
Proc <sub>11</sub>	40.00	C(9)	1279	14736	16898
Proc <sub>22</sub>	<b>82.00</b>	<b>C(20)</b>	<b>2557</b>	<b>17565</b>	<b>32420</b>



# Results compared to CPU

Comparing to CPU over GPU

- GPU is the standard approach for hardware acceleration.
- CPU is closer in TDP.
- CPU is more approachable and making a GPU version would require higher development time.
- already reached some board limitations.

MD5: 4.5 times faster than any comparable CPU version.

AES: Proximity of the C# version, but cannot compete AES-NI.

SHA: Only half the speed of i5, but faster than ARM processor. Potential for improvement.

ChaCha: Reaches the bandwidth limit of the Zynq board. Not quite i5 speed. Doubling processes: 2714 MBps. Only 150 more than the previous version.

MD5						
	Naive	Proc64	C#	C	OpenSSL <sub>low</sub>	OpenSSL <sub>high</sub>
Pi	152	3210	287	256	42	293
i5	152	3210	604	622	81	691

AES						
	Naive	Proc11	C#	C	OpenSSL <sub>low</sub>	OpenSSL <sub>high</sub>
Pi	400	1963	70	198	72	89
i5	400	1699	1963	340	847	5722

SHA					
	Naive	Proc4	C#	OpenSSL <sub>low</sub>	OpenSSL <sub>high</sub>
Pi	134	224	163	42	165
i5	134	224	438	61	461

ChaCha				
	Naive	Proc	OpenSSL <sub>low</sub>	OpenSSL <sub>high</sub>
Pi	80	2557	84	307
i5	80	2557	388	3092



# Power usage

Why use TDP?

- Selling point for FPGAs
- Only possibility in the current stage.
- Time consuming to do actual tests.

All of our versions is significantly more power efficient than the CPUs:

FPGA	1.765W
i5	65.000W
Pi	7.500W



# Future work

## Critical work

- The Dependency routing needs to be fixed in the hashing functions.
- Make hashes able to switch between messages to circumvent the stalling.
- Make a useful interface to expose our implementations.

## Optimizing work

- Investigate SHA hopefully getting the performance to reasonable levels
- Test if different approaches of AES approaches could yield better performance. Sugestions:
  1. Naive but pipelined
  2. Stateful BRAM.

## Comparing work

- Compare to other research papers results which often are written in HDL to see if SME can actually provide comparable results. This would require better a better FPGA.
- Test against a GPU.



Questions?

# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE





# HEHE



# HEHE



# HEHE



# HEHE



# HEHE



# HEHE

