



Faculty of Science


Cryptographic Library for FPGAs using SME

A Bachelor Project Defense

Jacob Herbst (mrw148) and
Jonas Flach-Jensen (sjm233)
Institute of Computer Science (DIKU)

June 22, 2021

Slide 1/14





Agenda

- Why FPGA's and SME?
- Presentation of MD5 and AES
- Results
- Suggestions for future work



Motivation - Why use FPGAs?

- Architecture
- Configurable not only in computation but also in interface
- Often fast as the overhead from generality is omitted
- Often lower power consumption than CPU's
- Difficult to program (Xilinx Vivado 🏠)



Motivation - Why use SME?

- Synchronous Message Exchange
- High-level software programming for Hardware
- Based on Communicating Sequential Processes
- Consists of; a simple process, and a clocked process with input and output busses, and an onTick() function.
Together with a simulation process for testing



MD5

- Cryptographic hash function
- Merkle-Damgaard construction
- Four different compression functions, 64 rounds

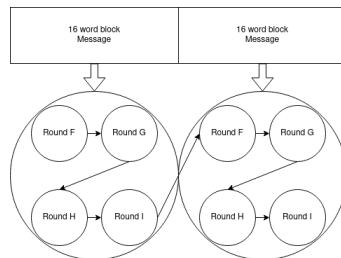


Figure: MD5 round




Figure: Merkle-Damgaard construction



MD5 - Optimizations

- Naive; 1 simple process, 4 busses
- Pipelined version; clocked process for preprocessing and each compression stage
- Same idea for SHA-2




Figure: MD5 pipeline



MD5 - Pipeline

- Stalls on large messages due to data dependency
- Solution is enabling multiple inputs

Independent message blocks									
clock	0	1	2	3	4	5	6	7	8
	P ₁	M ₁	F ₁	G ₁	H ₁	I ₁	C ₁		
		P ₂	M ₂	F ₂	G ₂	H ₂	I ₂	C ₂	

Dependent message blocks												
clock	0	1	2	3	4	5	6	7	8	9	10	11
	P ₁	M ₁	F ₁	G ₁	H ₁	I ₁	C ₁					
		P ₂	M ₂	-	-	-	-	F ₂	G ₂	H ₂	I ₂	C ₂



AES

- The algorithm is Rijndael. It is a Block Cipher and a Substitution-permutation (SP) network.
- Four steps, or lookups:

$$W_i \begin{cases} K_i & \text{if } i < N \\ W_{i-N} \oplus \text{SubWords}(W_{i-1} \lll 8) \oplus \text{rcon}_{i/N} & \text{if } i \geq N \text{ and } i \equiv 0 \pmod{N} \\ W_{i-N} \oplus \text{SubWords}(W_{i-1}) & \text{if } i \geq N, N > 6, \text{ and } i \equiv 4 \pmod{N} \\ W_{i-N} \oplus W_{i-1} & \text{otherwise} \end{cases}$$

$$T_0[a] = \begin{bmatrix} S[a] \cdot 02_{16} \\ S[a] \\ S[a] \\ S[a] \cdot 03_{16} \end{bmatrix} \quad T_1[a] = \begin{bmatrix} S[a] \cdot 03_{16} \\ S[a] \cdot 02_{16} \\ S[a] \\ S[a] \end{bmatrix} \quad T_2[a] = \begin{bmatrix} S[a] \\ S[a] \cdot 03_{16} \\ S[a] \cdot 02_{16} \\ S[a] \end{bmatrix} \quad T_3[a] = \begin{bmatrix} S[a] \\ S[a] \\ S[a] \cdot 03_{16} \\ S[a] \cdot 02_{16} \end{bmatrix}$$

$$e_j = T_0[a_{0,3}] \oplus T_1[a_{1,2}] \oplus T_2[a_{2,1}] \oplus T_3[a_{3,0}] \oplus W_j$$



AES - Optimization

- Fast naive version
- Pipelined by splitting up rounds
- no data dependency



Results from pipelining

- $hi(x) = x + 2 \cdot blocks$
- $lo(x) = 2 + x \cdot blocks$
- $C(x) = x + 2 \cdot blocks$
- Streaming AXI
- Zedboard bus: 2132 MBps

More processes is better throughput (to a certain extend)

MD5: Is easily optimised beyond the memory limit. 20x faster than Naive version.

SHA: Worst. Hard to optimize because of message expansion. 2x faster.

AES: Reached limit with current approach? Potentially other approaches can reach higher. 4x faster.

ChaCha: Starting to reach limit of board (100000 FF on board). 34x faster.

MD5							
Version	f_{max} (Mhz)	$clocks_{hi}$	$TP(MBps)_{hi}$	$clocks_{lo}$	$TP(MBps)_{lo}$	LUT	FF
Naive	2.38	b	152	b	152	11607	2304
Proc ₄	9.50	hi(6)	266	lo(6)	101	10247	5226
Proc ₈	19.00	hi(10)	532	lo(10)	122	10087	7538
Proc ₁₆	33.50	hi(18)	937	lo(18)	119	10206	12162
Proc ₃₂	65.00	hi(34)	1817	lo(34)	123	10149	21347
Proc₆₄	115.00	hi(66)	3209	lo(66)	112	10350	39718

SHA							
Version	f_{max} (Mhz)	$clocks_{hi}$	$TP(MBps)_{hi}$	$clocks_{lo}$	$TP(MBps)_{lo}$	LUT	FF
Naive	2.1	b	134.4	b	134.4	24330	2560
Proc₄	8.0	hi(6)	223.9	lo(6)	85.3	24466	8938
Proc ₈	8.0	hi(10)	223.8	lo(10)	51.2	24756	14066

AES						
Version	f_{max} (Mhz)	clocks	$TP(MBps)$	LUT	FF	BRAM
Naive	22	b	352	10612	3195	0
TBox	25	b	400	16458	3195	0
Proc ₄	68	C(3)	544	16474	2817	0
Proc ₁₁	208	C(12)	1663	15659	4383	0
Proc₂₂	217	C(24)	1735	15454	7401	0
BRAM ₁₁	195	C(31)	1556	10012	10398	72
Alt₃₆	240	C(38)	1916	6993	10808	40

ChaCha					
Version	f_{max} (Mhz)	clocks	$TP(MBps)$	LUT	FF
Naive	1.25	b	80	14670	3457
Proc ₁₁	40.00	C(9)	1279	14736	16898
Proc ₂₂	82.00	C(20)	2557	17565	32420
Proc₄₄	85.00	C(40)	2715	17612	62436



Results compared to CPU

Comparing to CPU over GPU

- GPU is the standard approach for hardware acceleration.
- CPU is more approachable and making a GPU version would require higher development time.
- already reached some board limitations.

MD5: ~4.5 times faster than any comparable CPU version.

AES: Proximity of the C# version, but cannot compete AES-NI.

SHA: Only half the speed of i5, but faster than ARM processor. Potential for improvement.

ChaCha: Percentagewise best but not quite speed of i5. Reaches the bandwidth limit of the Zynq board.

MD5						
	Naive	Proc _{i54}	C#	C	OpenSSL _{low}	OpenSSL _{high}
Pi	152	3210	287	256	42	293
i5	152	3210	604	622	81	691

AES						
	Naive	Proc _{i11}	C#	C	OpenSSL _{low}	OpenSSL _{high}
Pi	400	1916	70	198	72	89
i5	400	1916	1963	340	847	5722

SHA					
	Naive	Proc _{i4}	C#	OpenSSL _{low}	OpenSSL _{high}
Pi	134	224	163	42	165
i5	134	224	438	61	461

ChaCha				
	Naive	Proc	OpenSSL _{low}	OpenSSL _{high}
Pi	80	2715	84	307
i5	80	2715	388	3092



Power usage

Why use TDP?

- Selling point for FPGAs
- Only possibility in the current stage.
- Pretty diffuse concept as no standard way to measure it (Intel boost)

All of our versions seems significantly more power efficient than the CPUs:

FPGA		1.765W	
i5	65.0W	~6.500W	(i5 is 39 times bigger).
Pi	7.5W	~0.750W	



Future work

Critical work

- The Dependency routing needs to be fixed in the hashing functions.
- Make hashes able to switch between messages to circumvent the stalling.
- Make a useful interface to expose our implementations.

Optimizing work

- Investigate SHA hopefully getting the performance to reasonable levels
- Test if different approaches of AES approaches could yield better performance. Sugestions:
 1. Naive but pipelined
 2. Stateful BRAM.

Comparing work

- Compare to other research papers results which often are written in HDL to see if SME can actually provide comparable results. This would require better a better FPGA.
- Test against a GPU.



Questions?