



Faculty of Science


Cryptographic Library for FPGAs using SME

A Bachelor Project Defense

Jacob Herbst (mrw148) and
Jonas Flach-Jensen (sjm233)
Institute of Computer Science (DIKU)

June 18, 2021

Slide 1/23



Agenda

- Why is this project interesting?
- Presentation of MD5 and AES
- Results
- suggestions for future work



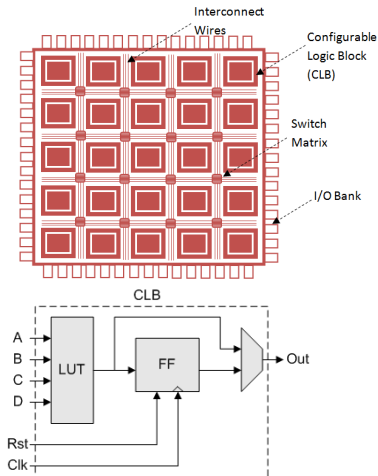
Motivation - Cryptography

- Is ubiquitous
- Often Hardware-focused
- This project has focused on implementing a variety of algorithms instead of nitpicking one algorithm.



Motivation - Why use FPGAs?

- Architecture
- Configurable not only in computation but also in interface
- Often fast as the overhead from generality is omitted
- Low power consumption (fixed precision)
- Difficult to program (Xilinx Vivado 🏠)



MD5

- Cryptographic hash function
- Merkle-Damgaard construction
- Four different compression functions, 64 rounds




Figure: MD5 round




Figure: Merkle-Damgaard construction

MD5 - Optimizations

- Naive; 1 simple process, 4 busses
- Pipelined version; clocked process for preprocessing and each compression stage
- Same idea for SHA-2




Figure: MD5 pipeline



MD5 - Pipeline

- Stalls on large messages
- Solution is enabling multiple inputs

Independent message blocks									
clock	0	1	2	3	4	5	6	7	8
	P_1	M_1	F_1	G_1	H_1	I_1	C_1		
		P_2	M_2	F_2	G_2	H_2	I_2	C_2	

Dependent message blocks												
clock	0	1	2	3	4	5	6	7	8	9	10	11
	P ₁	M ₁	F ₁	G ₁	H ₁	I ₁	C ₁					
		P ₂	M ₂	-	-	-	-	F ₂	G ₂	H ₂	I ₂	C ₂



AES



Results from pipelining

- $hi(x) = x + 2 \cdot blocks$
- $lo(x) = 2 + x \cdot blocks$
- $C(x) = x + 2 \cdot blocks$
- Zedboard bus: 2132 MBps

MD5: Is easily optimised beyond the memory limit.

SHA: Worst hard to optimize because of dependencies.

AES: Reached limit with current approach? Potentially other approaches can reach higher.

ChaCha: Easily reaches limit.

MD5							
Version	f_{max} (Mhz)	$clocks_{hi}$	$TP(MBps)_{hi}$	$clocks_{lo}$	$TP(MBps)_{lo}$	LUT	FF
Naïve	2.38	b	152	b	152	11607	2304
Proc ₄	9.50	hi(6)	266	lo(6)	101	10247	5226
Proc ₈	19.00	hi(10)	532	lo(10)	122	10087	7538
Proc ₁₆	33.50	hi(18)	937	lo(18)	119	10206	12162
Proc ₃₂	65.00	hi(34)	1817	lo(34)	123	10149	21347
Proc ₆₄	115.00	hi(66)	3209	lo(66)	112	10350	39718

SHA							
Version	f_{max} (Mhz)	$clocks_{hi}$	$TP(MBps)_{hi}$	$clocks_{lo}$	$TP(MBps)_{lo}$	LUT	FF
Naïve	2.1	b	134.4	b	134.4	24330	2560
Proc ₄	8.0	hi(6)	223.9	lo(6)	85.3	24466	8938
Proc ₈	8.0	hi(10)	223.8	lo(10)	51.2	24756	14066

AES						
Version	f_{max} (Mhz)	$clocks$	$TP(MBps)$	LUT	FF	BRAM
Naïve	22	b	352	10612	3195	0
TBox	25	b	400	16458	3195	0
Proc ₄	68	C(3)	544	16474	2817	0
Proc ₁₁	208	C(12)	1663	15659	4383	0
Proc ₂₂	217	C(24)	1662	15454	7401	0
BRAM ₁₁	195	C(31)	1556	10012	10398	72

ChaCha					
Version	f_{max} (Mhz)	$clocks$	$TP(MBps)$	LUT	FF
Naïve	1.25	b	80	14670	3457
Proc ₁₁	40.00	C(9)	1279	14736	16898
Proc ₂₂	82.00	C(20)	2557	17565	32420



Results compared to CPU

Comparing to CPU over GPU

- GPU is the standard approach for hardware acceleration.
- CPU is closer in TDP.
- CPU is more approachable and making a GPU version would require higher development time.
- already reached some board limitations.

MD5: 4.5 times faster than any comparable CPU version.

AES: Proximity of the C# version, but cannot compete AES-NI.

SHA: Only half the speed of i5, but faster than ARM processor. Potential for improvement.

ChaCha: Reaches the bandwidth limit of the Zynq board. Not quite i5 speed. Doubling processes: 2714 MBps. Only 150 more than the previous version.

MD5						
	Naive	Proc64	C#	C	OpenSSL _{low}	OpenSSL _{high}
Pi	152	3210	287	256	42	293
i5	152	3210	604	622	81	691

AES						
	Naive	Proc11	C#	C	OpenSSL _{low}	OpenSSL _{high}
Pi	400	1963	70	198	72	89
i5	400	1699	1963	340	847	5722

SHA					
	Naive	Proc4	C#	OpenSSL _{low}	OpenSSL _{high}
Pi	134	224	163	42	165
i5	134	224	438	61	461

ChaCha				
	Naive	Proc	OpenSSL _{low}	OpenSSL _{high}
Pi	80	2557	84	307
i5	80	2557	388	3092



Power usage

Why use TDP?

- Selling point for FPGAs
- Only possibility in the current stage.
- Time consuming to do actual tests.

All of our versions is significantly more power efficient than the CPUs:

FPGA	1.765W
i5	65.000W
Pi	7.500W



Future work

Critical work

- The Dependency routing needs to be fixed in the hashing functions.
- Make hashes able to switch between messages to circumvent the stalling.
- Make a useful interface to expose our implementations.

Optimizing work

- Investigate SHA hopefully getting the performance to reasonable levels
- Test if different approaches of AES approaches could yield better performance. Sugestions:
 1. Naive but pipelined
 2. Stateful BRAM.

Comparing work

- Compare to other research papers results which often are written in HDL to see if SME can actually provide comparable results. This would require better a better FPGA.
- Test against a GPU.



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE



HEHE

